

Subject: Notes on pandas library

Course: Artificial Intelligence - LM

Department: DISI (Department of Computer Science and Engineering) - University of Bologna, Italy

Author: Roberto Zanolli

Sources: https://pandas.pydata.org/docs/user_guide/index.html#user-guide

Pandas library for data science

Pandas stands for panel data.

It allows efficient management of excel like tables called **dataframes** and is mostly used for data science tasks.

To work with pandas of course we need to import it, it is a convention to alias the import as pd.

```
import pandas as pd
```

It introduces the most elementary data structure of the library: the **series**.

Series

A **series** is a one-dimensional labelled array.

```
fruits = pd.Series(['apple', 'banana', 'cherry', 'date', 'elderberry'],
                   index=['a', 'b', 'c', 'd', 'e'])

print(fruits)
print(f"\nIndex: {fruits.index.tolist()}")
```

Output:

```
a      apple
b     banana
c    cherry
d      date
e  elderberry
dtype: object

Index: ['a', 'b', 'c', 'd', 'e']
```

iloc method

```
print("== iloc examples (position-based) ==")
print(fruits.iloc[0])
```

```
print(fruits.iloc[2])
print(fruits.iloc[1:4]) # slice from position 1 to 3
```

Output:

```
== iloc examples (position-based) ==
apple
cherry
b    banana
c    cherry
d    date
dtype: object
```

loc method

```
print("\n== loc examples (label-based) ==")
print(fruits.loc['a'])
print(fruits.loc['c'])
print(fruits.loc['b':'d']) # slice from label 'b' to 'd' (inclusive!)
```

Output:

```
== loc examples (label-based) ==
apple
cherry
b    banana
c    cherry
d    date
dtype: object
```

Remember: `iloc` uses Python-style indexing (0-based, exclusive end), while `loc` uses label-based indexing (inclusive end).

Updating and filtering series

Here's a super easy example using days of the week and calories:

Creating the Series

```
calories = pd.Series([2200, 1800, 2100, 1900, 2300, 2800, 1700],
                     index=['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun'])

print("Original Calorie Data:")
print(calories)
```

Output:

```
Original Calorie Data:
Mon    2200
Tue    1800
```

```
Wed    2100  
Thu    1900  
Fri    2300  
Sat    2800  
Sun    1700  
dtype: int64
```

Filtering

```
print("\n==== FILTERING ===")  
  
# filter days with more than 2000 calories  
high_calorie_days = calories[calories > 2000]  
print("Days with >2000 calories:")  
print(high_calorie_days)  
  
# filter specific days using labels  
weekend_calories = calories[['Sat', 'Sun']]  
print("\nWeekend calories:")  
print(weekend_calories)  
  
# filter using multiple conditions  
medium_calories = calories[(calories >= 1800) & (calories <= 2200)]  
print("\nMedium calorie days (1800-2200):")  
print(medium_calories)
```

Output:

```
==== FILTERING ====  
Days with >2000 calories:  
Mon    2200  
Wed    2100  
Fri    2300  
Sat    2800  
dtype: int64  
  
Weekend calories:  
Sat    2800  
Sun    1700  
dtype: int64  
  
Medium calorie days (1800-2200):  
Mon    2200  
Tue    1800  
Wed    2100  
Thu    1900  
dtype: int64
```

updating

```
print("\n==== UPDATING VALUES ===")  
  
# change Monday's calories (position 0)  
calories.iloc[0] = 2000
```

```

print("After updating Monday (position 0):")
print(calories)

# change weekend calories (positions 5 and 6)
calories.iloc[5:7] = [2600, 1750]
print("\nAfter updating weekend (positions 5-6):")
print(calories)

# update every other day using positions
calories.iloc[1:6:2] = [1850, 1950, 2250] # update Tue, Thu, Fri
print("\nAfter updating Tue, Thu, Fri (positions 1,3,5):")
print(calories)

```

Output:

```

== UPDATING VALUES ==
After updating Monday (position 0):
Mon    2000
Tue    1800
Wed    2100
Thu    1900
Fri    2300
Sat    2800
Sun    1700
dtype: int64

After updating weekend (positions 5-6):
Mon    2000
Tue    1800
Wed    2100
Thu    1900
Fri    2300
Sat    2600
Sun    1750
dtype: int64

After updating Tue, Thu, Fri (positions 1,3,5):
Mon    2000
Tue    1850
Wed    2100
Thu    1950
Fri    2250
Sat    2600
Sun    1750
dtype: int64

```

Dataframes

A DataFrame is a **two-dimensional labelled data structure** with columns of potentially different types.

```

# create a DataFrame from a dictionary
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'Diana'],

```

```

'Age': [25, 30, 35, 28],
'City': ['New York', 'London', 'Tokyo', 'Paris'],
'Salary': [50000, 60000, 70000, 55000]
}

df = pd.DataFrame(data)
print("Original DataFrame:")
print(df)

```

Output:

```

Original DataFrame:
   Name  Age      City  Salary
0  Alice  25  New York   50000
1    Bob  30    London   60000
2 Charlie  35     Tokyo   70000
3 Diana  28    Paris   55000

```

It is possible to access to dataframes cells by using the notation [row,column].

iloc method

```

print("== iloc examples (position-based) ==")
print(df.iloc[0])          # first row
print(df.iloc[2, 1])        # element at row 2, column 1 (Age=35)
print(df.iloc[1:3])         # rows 1 to 2 (exclusive end)

```

Output:

```

== iloc examples (position-based) ==
Name      Alice
Age       25
City    New York
Salary    50000
Name: 0, dtype: object
35
   Name  Age      City  Salary
0  Alice  25  New York   50000
1    Bob  30    London   60000
2 Charlie  35     Tokyo   70000

```

loc method

```

print("\n== loc examples (label-based) ==")
print(df.loc[0])          # row with index label 0
print(df.loc[2, 'City'])    # element at row 2, column 'City'
print(df.loc[1:3])         # rows 1 to 3 (inclusive end)

```

Output:

```

== loc examples (label-based) ==
Name      Alice
Age       25

```

```

City      New York
Salary     50000
Name: 0, dtype: object
Tokyo
      Name  Age   City  Salary
1      Bob   30  London   60000
2  Charlie   35  Tokyo   70000
3   Diana   28  Paris   55000

```

The methods loc and iloc works the same as in series

Filtering and updating DataFrames

Filtering

```

print("\n==== FILTERING ===")

# filter people over 28 years old
older_than_28 = df[df['Age'] > 28]
print("People older than 28:")
print(older_than_28)

# filter specific columns
name_city = df[['Name', 'City']]
print("\nOnly Name and City columns:")
print(name_city)

# multiple conditions
high_earners_young = df[(df['Salary'] > 55000) & (df['Age'] < 35)]
print("\nHigh earners under 35:")
print(high_earners_young)

```

Output:

```

==== FILTERING ====
People older than 28:
      Name  Age   City  Salary
1      Bob   30  London   60000
2  Charlie   35  Tokyo   70000

Only Name and City columns:
      Name      City
0    Alice  New York
1      Bob    London
2  Charlie    Tokyo
3   Diana    Paris

High earners under 35:
      Name  Age   City  Salary
1      Bob   30  London   60000

```

Updating values

```

print("\n==== UPDATING VALUES ===")

# update single value - Alice's salary (row 0, column 'Salary')
df.loc[0, 'Salary'] = 52000
print("After updating Alice's salary:")
print(df)

# update multiple values - give raises to rows 1 and 3
df.loc[[1, 3], 'Salary'] = [65000, 58000]
print("\nAfter giving raises to Bob and Diana:")
print(df)

# update using iloc - change cities for first two people
df.iloc[0:2, 2] = ['Boston', 'Manchester']
print("\nAfter updating cities (positions 0-1):")
print(df)

```

Output:

```

==== UPDATING VALUES ===
After updating Alice's salary:
      Name  Age      City  Salary
0    Alice   25  New York   52000
1      Bob   30    London   60000
2  Charlie   35     Tokyo   70000
3   Diana   28    Paris   55000

After giving raises to Bob and Diana:
      Name  Age      City  Salary
0    Alice   25  New York   52000
1      Bob   30    London   65000
2  Charlie   35     Tokyo   70000
3   Diana   28    Paris   58000

After updating cities (positions 0-1):
      Name  Age      City  Salary
0    Alice   25     Boston   52000
1      Bob   30  Manchester   65000
2  Charlie   35     Tokyo   70000
3   Diana   28    Paris   58000

```

Adding a Row using concat

Adding a single row using concat

```

print("\n==== ADDING A ROW WITH CONCAT ===")

# create a new row as a DataFrame
new_person = pd.DataFrame({
    'Name': ['Evan'],
    'Age': [32],
    'City': ['Berlin'],
    'Salary': [62000]
})

```

```
# use concat to add the new row
df = pd.concat([df, new_person], ignore_index=True)
print("After adding Evan:")
print(df)
```

Output:

```
==== ADDING A ROW WITH CONCAT ====
After adding Evan:
      Name  Age     City   Salary
0    Alice  25  New York  50000
1     Bob  30    London  60000
2  Charlie  35    Tokyo  70000
3   Diana  28    Paris  55000
4    Evan  32  Berlin  62000
```

Adding multiple rows

```
print("\n==== ADDING MULTIPLE ROWS ===")

# create multiple new rows
new_people = pd.DataFrame({
    'Name': ['Fiona', 'George'],
    'Age': [29, 40],
    'City': ['Sydney', 'Toronto'],
    'Salary': [58000, 75000]
})

df = pd.concat([df, new_people], ignore_index=True)
print("After adding Fiona and George:")
print(df)
```

Output:

```
==== ADDING MULTIPLE ROWS ====
After adding Fiona and George:
      Name  Age     City   Salary
0    Alice  25  New York  50000
1     Bob  30    London  60000
2  Charlie  35    Tokyo  70000
3   Diana  28    Paris  55000
4    Evan  32  Berlin  62000
5    Fiona  29    Sydney  58000
6    George  40  Toronto  75000
```

Remember: Always use `ignore_index=True` when adding rows with `concat()` to maintain clean sequential indexing.

Important: Using `ignore_index`

```
print("\n==== WITHOUT ignore_index ===")

# show what happens without ignore_index
```

```
temp_df = pd.concat([df, new_person])
print("Without ignore_index (duplicate indices):")
print(temp_df.tail(3))
```

Output:

```
== WITHOUT ignore_index ==
Without ignore_index (duplicate indices):
   Name  Age  City  Salary
6  George  40  Toronto  75000
0    Evan  32  Berlin  62000
```

If you had a custom index in the first `df` you are adding rows to, you can put a specific index using `index = ["placeholder"]` also in the second `df`.

Alternative: Adding row as Series

```
print("\n== ADDING ROW AS SERIES ==")

# create new row as a Series
new_row = pd.Series({
    'Name': 'Hannah',
    'Age': 27,
    'City': 'Rome',
    'Salary': 53000
})

# convert to DataFrame and concat
df = pd.concat([df, pd.DataFrame([new_row])], ignore_index=True)
print("After adding Hannah:")
print(df.tail(3))
```

Output:

```
== ADDING ROW AS SERIES ==
After adding Hannah:
   Name  Age  City  Salary
5  Fiona  29  Sydney  58000
6  George  40  Toronto  75000
7  Hannah  27    Rome  53000
```

Key points

- `concat()` combines DataFrames along a particular axis (rows by default)
- `ignore_index=True` resets the index to avoid duplicate indexes
- New rows must be passed as **DataFrame** (even if single row)
- Alternative to `append()` (which is deprecated)

Reading data from file

Pandas allows us to read csv and excel files and load them directly inside a dataframe.

It also allows to set custom **separators** and **index**.

```
# basic csv reading
df = pd.read_csv('data.csv')
print("basic csv reading:")
print(df.head())
```

Output:

```
basic csv reading:
   id    name  age      city
0   1    alice  25  new york
1   2      bob  30    london
2   3  charlie  35     tokyo
```

The main methods to interact with the filesystem are:

- **read_csv()** - reads data from comma-separated values file into a dataframe.
- **read_excel()** - reads data from excel spreadsheet into a dataframe.
- **to_csv()** - writes dataframe to a csv file.
- **to_excel()** - writes dataframe to an excel file.

reading with custom settings

```
# reading csv with custom index and delimiter
df = pd.read_csv('data.csv', index_col='id', delimiter=',')
print("\nwith custom index (id column):")
print(df.head())
```

Output:

```
with custom index (id column):
      name  age      city
id
1    alice  25  new york
2      bob  30    london
3  charlie  35     tokyo
```

handling different delimiters

```
# for tab-separated files
df_tsv = pd.read_csv('data.tsv', delimiter='\t')

# for semicolon-separated files
df_semicolon = pd.read_csv('data.csv', delimiter=';')

# for files with custom separators
df_custom = pd.read_csv('data.txt', delimiter='|')
```

using head() method

```
# head() shows first few rows (default 5)
print("first 3 rows:")
print(df.head(3))

print("\nfirst 5 rows (default):")
print(df.head())
```

Output:

```
first 3 rows:
      name    age      city
id
1     alice   25  new york
2       bob   30    london
3  charlie   35     tokyo

first 5 rows (default):
      name    age      city
id
1     alice   25  new york
2       bob   30    london
3  charlie   35     tokyo
```

The method **tail()** does the opposite.

using **describe()** method

```
# describe() shows statistical summary for numerical columns
print("statistical summary:")
print(df.describe())
```

Output:

```
statistical summary:
      age
count  3.000000
mean  30.000000
std   5.000000
min   25.000000
25%   27.500000
50%   30.000000
75%   32.500000
max   35.000000
```

complete example with all methods

```
# complete file reading example
df = pd.read_csv('employees.csv',
                  index_col='employee_id', # set custom index
                  delimiter=',')          # set delimiter

print("dataset overview:")
print(f"shape: {df.shape}") # shows (rows, columns)
```

```

print(f"columns: {df.columns.tolist()}")

print("\nfirst 5 rows:")
print(df.head())

print("\nstatistical summary:")
print(df.describe())

#print("\nfirst 10 rows:")
#print(df.head(10))

```

Output:

```

dataset overview:
shape: (100, 4)
columns: ['name', 'age', 'department', 'salary']

first 5 rows:
      name  age department  salary
employee_id
1       alice  25        hr  50000
2         bob  30      sales  60000
3     charlie  35  engineering  70000
4       diana  28  marketing  55000
5       evan  32         it  62000

statistical summary:
      age      salary
count 100.000000 100.000000
mean   35.200000 65230.000000
std    8.503872 14923.654321
min   22.000000 40000.000000
25%  29.000000 55000.000000
50%  35.000000 65000.000000
75%  41.000000 75000.000000
max   55.000000 95000.000000

```

Main data inspection methods and attributes

The essential pandas methods for inspecting the data are:

- head()** - shows the first few rows to quickly preview your data
- tail()** - displays the last few rows to see the end of your dataset
- info()** - provides technical overview including data types and missing values
- describe()** - gives statistical summary like mean, min, max for numerical columns

There are also attributes of dataframes we can access (not methods):

- shape** - tells you the size of your dataset in (rows, columns) format
- columns** - lists all the column names in your dataframe
- index** - shows the row labels or identifiers used in your data

Examples:

```

df = pd.DataFrame({
    'name': ['alice', 'bob', 'charlie', 'diana'],
    'age': [25, 30, 35, 28],
    'city': ['new york', 'london', 'tokyo', 'paris'],
    'salary': [50000, 60000, 70000, 55000]
}, index=['emp001', 'emp002', 'emp003', 'emp004'])

print("our dataframe:")
print(df)

```

Output:

```

our dataframe:
      name  age      city  salary
emp001   alice  25  new york   50000
emp002     bob  30    london   60000
emp003  charlie  35     tokyo   70000
emp004   diana  28     paris   55000

```

using shape

```

print(f"dataframe shape: {df.shape}")
print(f"we have {df.shape[0]} rows and {df.shape[1]} columns")

```

Output:

```

dataframe shape: (4, 4)
we have 4 rows and 4 columns

```

using columns

```

print(f"columns: {df.columns.tolist()}")
print(f"number of columns: {len(df.columns)}")
print("individual columns:")
for col in df.columns:
    print(f" - {col}")

```

Output:

```

columns: ['name', 'age', 'city', 'salary']
number of columns: 4
individual columns:
 - name
 - age
 - city
 - salary

```

using index

```
print(f"index: {df.index.tolist()}")
print(f"index type: {type(df.index)}")
print("row identifiers:")
for idx in df.index:
    print(f" - {idx}")
```

Output:

```
index: ['emp001', 'emp002', 'emp003', 'emp004']
index type: <class 'pandas.core.indexes.base.Index'>
row identifiers:
 - emp001
 - emp002
 - emp003
 - emp004
```

Data analysis methods

When working with data analysis in pandas, aggregation methods are essential for summarizing and understanding your dataset. These functions allow to transform raw data into insights by calculating statistics, counting occurrences, and grouping information logically.

using groupby

```
# Group data by person to analyze each individual separately
person_groups = calories.groupby('person')
print("Data grouped by person for individual analysis")
```

using mean

```
# Calculate average calories per person
avg_calories = calories.groupby('person')['calories'].mean()
print("Average daily calories per person:")
print(avg_calories)
```

Output:

```
Average daily calories per person:
person
Alice    2114.29
Bob      2366.67
Name: calories, dtype: float64
```

using sum

```
# Calculate total calories consumed per person
total_calories = calories.groupby('person')['calories'].sum()
print("Total calories consumed per person:")
print(total_calories)
```

Output:

```
Total calories consumed per person:  
person  
Alice    14800  
Bob      7100  
Name: calories, dtype: int64
```

using count

```
# Count number of days tracked per person  
days_tracked = calories.groupby('person')['day'].count()  
print("Number of days tracked per person:")  
print(days_tracked)
```

Output:

```
Number of days tracked per person:  
person  
Alice    7  
Bob      3  
Name: day, dtype: int64
```

using agg

```
# Comprehensive analysis using multiple aggregation functions  
summary = calories.groupby('person').agg({  
    'calories': ['mean', 'max', 'min', 'sum'],  
    'meals': ['mean', 'count']  
})  
print("Complete nutritional summary:")  
print(summary)
```

Output:

```
Complete nutritional summary:  
          calories               meals  
              mean     max   min   sum   mean  count  
person  
Alice  2114.29  2800  1700  14800  3.00     7  
Bob    2366.67  2500  2200   7100  3.67     3
```

Adding columns

A DataFrame can be extended by adding new labelled columns, created from constants, lists, operations or functions.

adding a column with a constant value

```
df['Country'] = 'Unknown'
```

```
print(df)
```

Output:

```
Name    Age     City   Salary  Country  
0  Alice    25  New York  50000  Unknown  
1    Bob    30    London  60000  Unknown  
2 Charlie   35    Tokyo  70000  Unknown  
3 Diana    28    Paris  55000  Unknown
```

adding a column from a list

```
df['Experience'] = [1, 4, 7, 3]  
print(df)
```

Output:

```
Name    Age     City   Salary  Country  Experience  
0  Alice    25  New York  50000  Unknown      1  
1    Bob    30    London  60000  Unknown      4  
2 Charlie   35    Tokyo  70000  Unknown      7  
3 Diana    28    Paris  55000  Unknown      3
```

adding a column using operations on existing columns

```
df['Salary_in_k'] = df['Salary'] / 1000  
print(df[['Name', 'Salary', 'Salary_in_k']])
```

Output:

```
Name    Salary  Salary_in_k  
0  Alice    50000      50.0  
1    Bob    60000      60.0  
2 Charlie   70000      70.0  
3 Diana    55000      55.0
```

adding a column with conditional logic

```
df['Is_Adult'] = df['Age'] >= 18  
print(df[['Name', 'Age', 'Is_Adult']])
```

Output:

```
Name    Age  Is_Adult  
0  Alice    25      True  
1    Bob    30      True  
2 Charlie   35      True  
3 Diana    28      True
```

adding a column using apply() and a custom function

```
def categorize_salary(s):
    return "High" if s > 60000 else "Normal"

df['Salary_Level'] = df['Salary'].apply(categorize_salary)
print(df[['Name', 'Salary', 'Salary_Level']])
```

Output:

```
Name  Salary  Salary_Level
0   Alice    50000      Normal
1     Bob    60000      Normal
2  Charlie   70000      High
3   Diana   55000      Normal
```

adding multiple columns with assign()

```
df = df.assign(
    Age_plus_10 = df['Age'] + 10,
    Double_Salary = df['Salary'] * 2
)

print(df)
```

Output:

```
Name  Age      City  Salary  Country  Experience  Salary_in_k  Is_Adult
Salary_Level  Age_plus_10  Double_Salary
0   Alice    25  New York   50000  Unknown       1        50.0    True
Normal          35            100000
1     Bob    30    London   60000  Unknown       4        60.0    True
Normal          40            120000
2  Charlie   35     Tokyo   70000  Unknown       7        70.0    True
High            45            140000
3   Diana    28    Paris   55000  Unknown       3        55.0    True
Normal          38            110000
```

Removing columns

You can remove one or more columns using the drop() method.

Remember: by default, drop returns a new DataFrame, unless inplace=True is specified.

removing a single column

```
df_no_country = df.drop('Country', axis=1)
print(df_no_country)
```

Output:

	Name	Age	City	Salary	Experience	Salary_in_k	Is_Adult	Salary_Level
Age_plus_10	Double_Salary							
0	Alice	25	New York	50000	1	50.0	True	Normal
35		100000						
1	Bob	30	London	60000	4	60.0	True	Normal
40		120000						
2	Charlie	35	Tokyo	70000	7	70.0	True	High
45		140000						
3	Diana	28	Paris	55000	3	55.0	True	Normal
38		110000						

removing multiple columns

```
df_no_salary = df.drop(['Salary', 'Salary_in_k'], axis=1)
print(df_no_salary)
```

Output:

	Name	Age	City	Country	Experience	Is_Adult	Salary_Level	Age_plus_10
Double_Salary								
0	Alice	25	New York	Unknown	1	True	Normal	35
100000								
1	Bob	30	London	Unknown	4	True	Normal	40
120000								
2	Charlie	35	Tokyo	Unknown	7	True	High	45
140000								
3	Diana	28	Paris	Unknown	3	True	Normal	38
110000								

removing a column permanently (inplace)

```
df.drop('Experience', axis=1, inplace=True)
print(df)
```

Output:

	Name	Age	City	Salary	Country	Salary_in_k	Is_Adult	Salary_Level
Age_plus_10	Double_Salary							
0	Alice	25	New York	50000	Unknown	50.0	True	Normal
35		100000						
1	Bob	30	London	60000	Unknown	60.0	True	Normal
40		120000						
2	Charlie	35	Tokyo	70000	Unknown	70.0	True	High
45		140000						
3	Diana	28	Paris	55000	Unknown	55.0	True	Normal
38		110000						

removing a column using del

```
del df['Country']
print(df)
```

Output:

	Name	Age	City	Salary	Salary_in_k	Is_Adult	Salary_Level	Age_plus_10
0	Alice	25	New York	50000	50.0	True	Normal	35
1	Bob	30	London	60000	60.0	True	Normal	40
2	Charlie	35	Tokyo	70000	70.0	True	High	45
3	Diana	28	Paris	55000	55.0	True	Normal	38

splitting X and y

In many machine learning workflows you need to separate the **features** (X) from the **target** (y).

Usually, X contains all columns except the target, and y contains only the target column.

example dataset

```
data = {
    'Age': [25, 30, 35, 28],
    'Salary': [50000, 60000, 70000, 55000],
    'City': ['New York', 'London', 'Tokyo', 'Paris'],
    'Purchased': [0, 1, 0, 1] # target column
}

df = pd.DataFrame(data)
print(df)
```

Output:

	Age	Salary	City	Purchased
0	25	50000	New York	0
1	30	60000	London	1
2	35	70000	Tokyo	0
3	28	55000	Paris	1

split X and y (classic approach)

```
X = df.drop('Purchased', axis=1)
y = df['Purchased']

print("X (features):")
print(X)

print("\ny (target):")
print(y)
```

Output:

```
X (features):
   Age  Salary      City
0   25  50000  New York
1   30  60000    London
2   35  70000    Tokyo
3   28  55000    Paris

y (target):
0    0
1    1
2    0
3    1
Name: Purchased, dtype: int64
```

alternative: selecting only feature columns

```
feature_cols = ['Age', 'Salary', 'City']
X = df[feature_cols]
y = df['Purchased']

print(X)
print(y)
```

Output:

```
   Age  Salary      City
0   25  50000  New York
1   30  60000    London
2   35  70000    Tokyo
3   28  55000    Paris

0    0
1    1
2    0
3    1
Name: Purchased, dtype: int64
```